

```
In [1]: import pandas as pd
df=pd.read_csv("diabetes.csv",header=None,names=["Pregnancies","Glucose","
```

```
In [2]: df.sample(2)
```

Out[2]:

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | Diabetes | Age | Cl |
|-----|-------------|---------|---------------|---------------|---------|------|----------|-----|----|
| 734 | 2 | 105 | 75 | 0 | 0 | 23.3 | 0.560 | 53 | |
| 590 | 11 | 111 | 84 | 40 | 0 | 46.8 | 0.925 | 45 | |

```
In [3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Pregnancies     768 non-null   int64
1   Glucose         768 non-null   int64
2   BloodPressure   768 non-null   int64
3   SkinThickness   768 non-null   int64
4   Insulin         768 non-null   int64
5   BMI             768 non-null   float64
6   Diabetes        768 non-null   float64
7   Age             768 non-null   int64
8   Class           768 non-null   int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
In [4]: x=df.iloc[:, :-1]
x
```

Out[4]:

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | Diabetes | Age |
|-----|-------------|---------|---------------|---------------|---------|------|----------|-----|
| 0 | 6 | 148 | 72 | 35 | 0 | 33.6 | 0.627 | 50 |
| 1 | 1 | 85 | 66 | 29 | 0 | 26.6 | 0.351 | 31 |
| 2 | 8 | 183 | 64 | 0 | 0 | 23.3 | 0.672 | 32 |
| 3 | 1 | 89 | 66 | 23 | 94 | 28.1 | 0.167 | 21 |
| 4 | 0 | 137 | 40 | 35 | 168 | 43.1 | 2.288 | 33 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 763 | 10 | 101 | 76 | 48 | 180 | 32.9 | 0.171 | 63 |
| 764 | 2 | 122 | 70 | 27 | 0 | 36.8 | 0.340 | 27 |
| 765 | 5 | 121 | 72 | 23 | 112 | 26.2 | 0.245 | 30 |
| 766 | 1 | 126 | 60 | 0 | 0 | 30.1 | 0.349 | 47 |
| 767 | 1 | 93 | 70 | 31 | 0 | 30.4 | 0.315 | 23 |

768 rows × 8 columns

```
In [16]: from sklearn.preprocessing import StandardScaler
```

```
In [17]: sc=StandardScaler()
xscaled=sc.fit_transform(x)
```

```
In [18]: x.shape
```

Out[18]: (768, 8)

```
In [19]: y=df.iloc[:, -1]
y
```

Out[19]:

| | |
|-----|----|
| 0 | 1 |
| 1 | 0 |
| 2 | 1 |
| 3 | 0 |
| 4 | 1 |
| ... | .. |
| 763 | 0 |
| 764 | 0 |
| 765 | 0 |
| 766 | 1 |
| 767 | 0 |

Name: Class, Length: 768, dtype: int64

```
In [20]: ▶ y.shape
```

```
Out[20]: (768,)
```

```
In [21]: ▶ from sklearn.model_selection import train_test_split
```

```
In [23]: ▶ xtrain,xtest,ytrain,ytest=train_test_split(x,y,random_state=1,test_size=0.
```

```
In [24]: ▶ xstrain,xstest,ystrain,ystest=train_test_split(x,y,random_state=1,test_siz
```

```
In [25]: ▶ xtrain.shape
```

```
Out[25]: (576, 8)
```

```
In [26]: ▶ xtest.shape
```

```
Out[26]: (192, 8)
```

```
In [27]: ▶ ytrain.shape
```

```
Out[27]: (576,)
```

```
In [28]: ▶ ytest.shape
```

```
Out[28]: (192,)
```

```
In [29]: ▶ from sklearn.linear_model import LogisticRegression
```

```
In [30]: ▶ model=LogisticRegression()
```

```
In [31]: ▶ model2=LogisticRegression()
```

```
In [32]: model.fit(xtrain,ytrain)
```

```
C:\Users\SAM\anaconda3\Lib\site-packages\sklearn\linear_model\_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
Out[32]: LogisticRegression
LogisticRegression()
```

```
In [33]: model2.fit(xtrain,ytrain)
```

```
C:\Users\SAM\anaconda3\Lib\site-packages\sklearn\linear_model\_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
Out[33]: LogisticRegression
LogisticRegression()
```

```
In [35]: df["Class"].value_counts()
```

```
Out[35]: Class
0      500
1      268
Name: count, dtype: int64
```

```
In [36]: predictions=model.predict(xtest)
```

In [37]: `predictions2=model2.predict(xtest)`

In [38]: `from sklearn.metrics import accuracy_score, confusion_matrix, classification`

In [39]: `accuracy_score(ytest, predictions)`

Out[39]: 0.7760416666666666

In [40]: `accuracy_score(ytest, predictions2)`

Out[40]: 0.7760416666666666

In [41]: `confusion_matrix(ytest, predictions)`

Out[41]: array([[109, 14],
[29, 40]], dtype=int64)

In [42]: `confusion_matrix(ytest, predictions2)`

Out[42]: array([[109, 14],
[29, 40]], dtype=int64)

In [44]: `ytest.value_counts()`

Out[44]: Class
0 123
1 69
Name: count, dtype: int64

In [45]: `print(classification_report(ytest, predictions))`

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.79 | 0.89 | 0.84 | 123 |
| 1 | 0.74 | 0.58 | 0.65 | 69 |
| accuracy | | | 0.78 | 192 |
| macro avg | 0.77 | 0.73 | 0.74 | 192 |
| weighted avg | 0.77 | 0.78 | 0.77 | 192 |

In [46]: `print(classification_report(ystest,predictions2))`

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.79 | 0.89 | 0.84 | 123 |
| 1 | 0.74 | 0.58 | 0.65 | 69 |
| accuracy | | | 0.78 | 192 |
| macro avg | 0.77 | 0.73 | 0.74 | 192 |
| weighted avg | 0.77 | 0.78 | 0.77 | 192 |

In []: